

Title: Streaming MapReduce & Hive using AWS
Student: Prateek Parasher
Lecturer: Edwina Sweeney
Course: M.Sc. in Computing in Big Data Analytics
Module: Big Data Architecture
Submission: 21st may 2018 via Blackboard

Introduction

Cloud computing is providing new power to industries all around the world and giving on demand computational power, storage, ram, different applications & other resources. This technology reduces the hardware configure and reduce maximum time in configuration like MapReduce, Hadoop this big data technology having very complicated configuration process & we have to make sure that our system having Linux / mac-os and compatible hardware to configure big data technologies & if we want to train heavily dataset then also we need high computational power

1. Amazon web services(AWS) is world leader in cloud computing services. In this project I am using aws providing lots of service in this project I am using few of them like S3 , EMR , EC2 so and so on

prior start my project details I would like give brief idea about aws , aws is currently providing their services nearly in 190 countries and they are keep expanding 42 zones for aws dealing with millions of active users

2. EMR stands for elastic MapReduce & emr is one of the aws services major component in emr is nodes, cluster & instances In amazon emr we can use Hadoop, spark components and usually in emr cluster having different nodes like master node, current node, task node.

➔ master node:- this is the main node that has all powers of controlling master node manages the running software compounds & manage the coordination between other tasks

Core node: - a slave node which is controlled by master node & slave node store all the data which run by other software component into the hdfs on the cluster

Task node:- this is also slave node but only responsible for running the tasks.

After getting overview of cluster & nodes we will talk about advantages of EMR

- ➔ Cost saving:- pay as you go depends upon the services we are using we have pay only for that like if we launch Ec2 using instance in west orgen region. We will be charged only that according to usages
- ➔ Aws integration :- for example if we using EMR service in our cluster we can integrate storage like AWS (S3) , amazon cloud watch to monitor the performance of cluster deployment, scalability, security etc

Amazon EMR is great service provider and make all difficult setup very easy, big data in different type of applications including machine learning, web indexing, bio informatics, log analysis, amazon EMR using Hadoop open source, which distribute data & processing across a resizable cluster Ec2(AWS)

Hadoop using distribution processing which is called MapReduce in which a mapper converts raw source data into key/value pairs and the find result comes they reduce the output to single set

In this project Initially, I created S3 bucket which is in AWS, S3 bucket for storage and creating s3 is very easy

Topology design

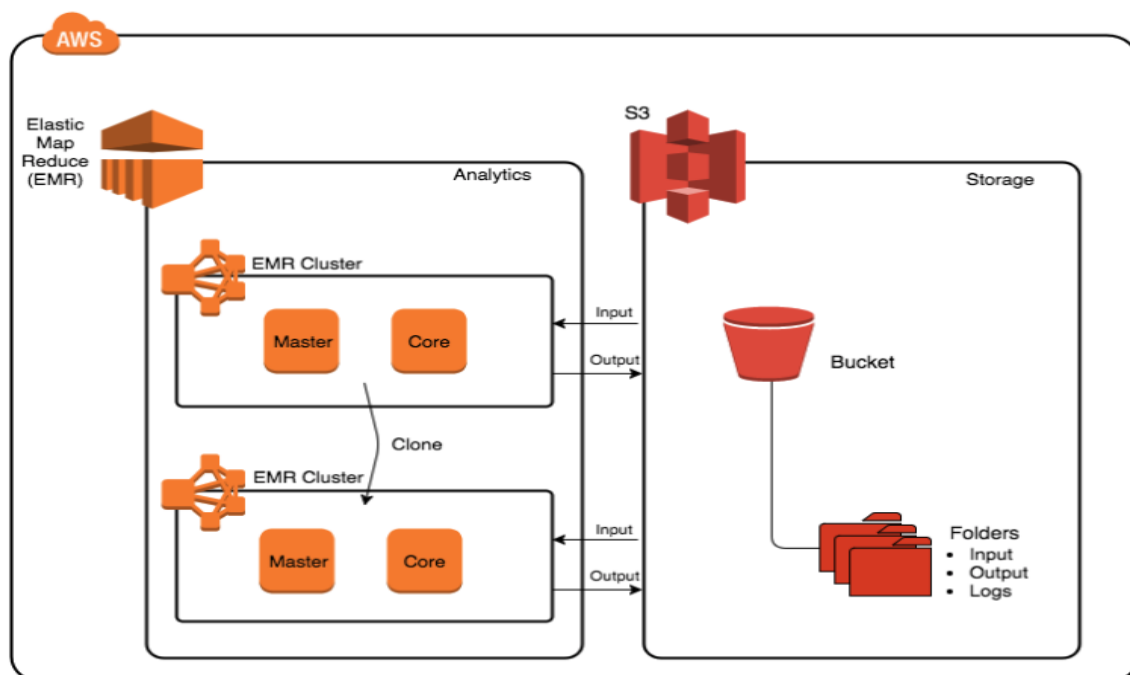


fig1 EMR in aws using S3 bucket

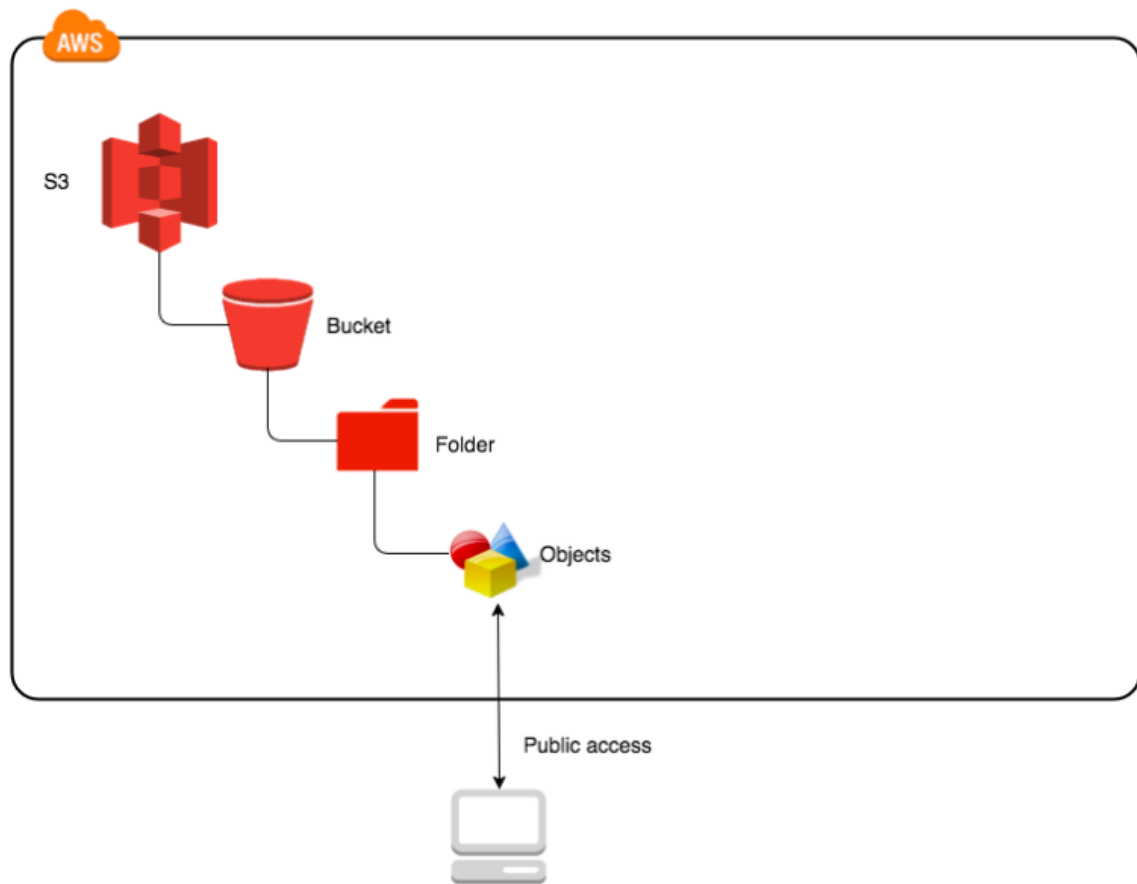


Fig 2 S3 bucket store system

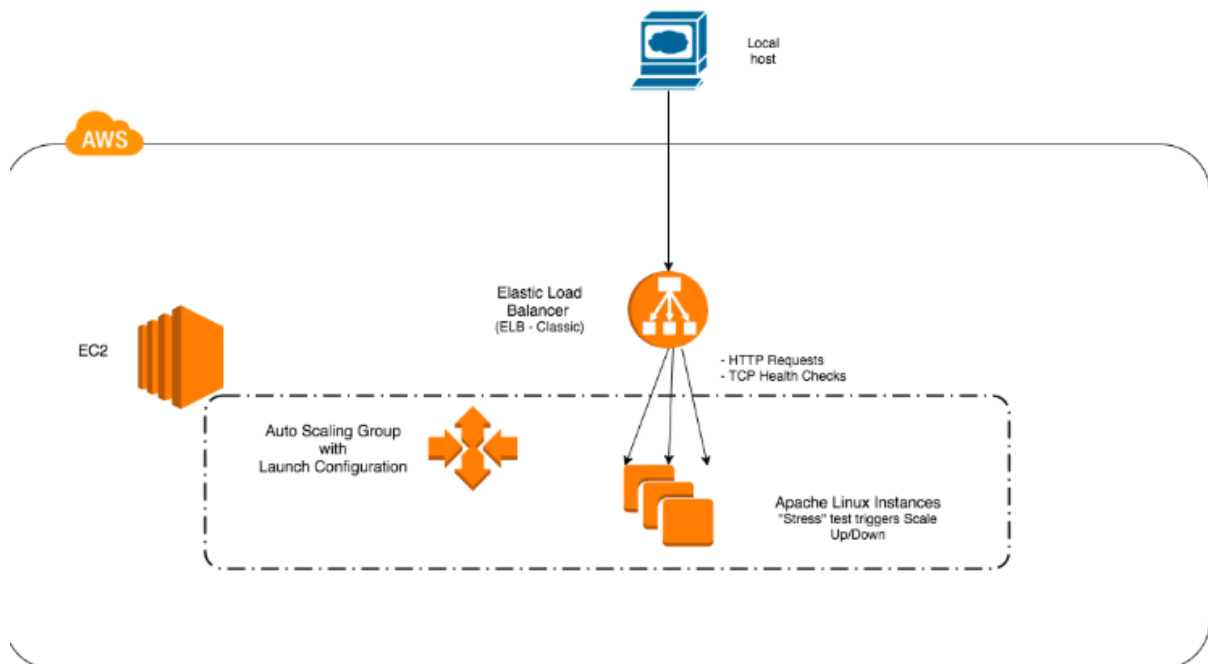


Fig 3 Ec2 cluster instance

Use cases

Zillow

3. Zillow is one of the largest real estate company in united states which delay with millions of home valuation data all the operation Zillow done in real-time using machine-learning calculation performance Zillow uses amazon kinesis streams input different type of data, sales & marketing transaction. they using machine learning algorithm all data store in amazon S3(aws) increasing in speed and accuracy using spark on emr (AWS), main benefits are prior using aws Zillow used to spend 1 day for high performance & accuracy now they are doing same process in 1 hr using AWS, they can easily scale up & compute capacity on demand

YELP

4. Yelp is more that decade year old company mainly focused on connecting people for businesses & help local business also yelp is backed by the community yelp dealing with 29 countries & 120 market through this expansion yelp faced few different challenges like storage , scaling up to overcome this problem they start using aws cloud computing services they start using single local instance of Hadoop and later on yelp move to EMR and yelp start using S3 for store the daily logs , picture , text , which is approx. 1.2 TB per day after consuming aws services yelp nearly save \$55,000 cost of hardware running per day.

Expedia

5. Expedia is one of the largest company in travel industry. Expedia company deals in hotels, travels & other facilities they are providing their services to 60 countries. They using aws amazon emr & amazon S3 for storage before that Expedia facing lot's of different problems like scaling up their business and Expedia approx. processes 240 request per second for security of the transaction in Expedia they using IAM aws security services to manage their security

Deployment Steps

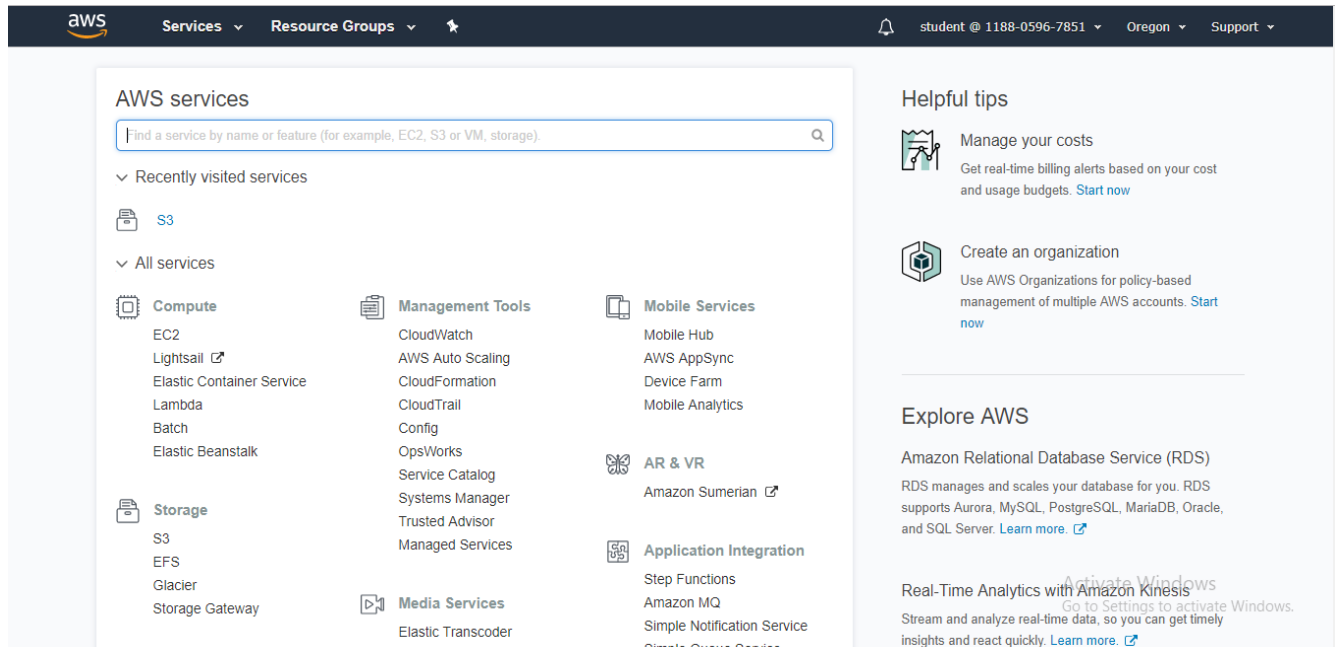


Fig 4 front view aws all services

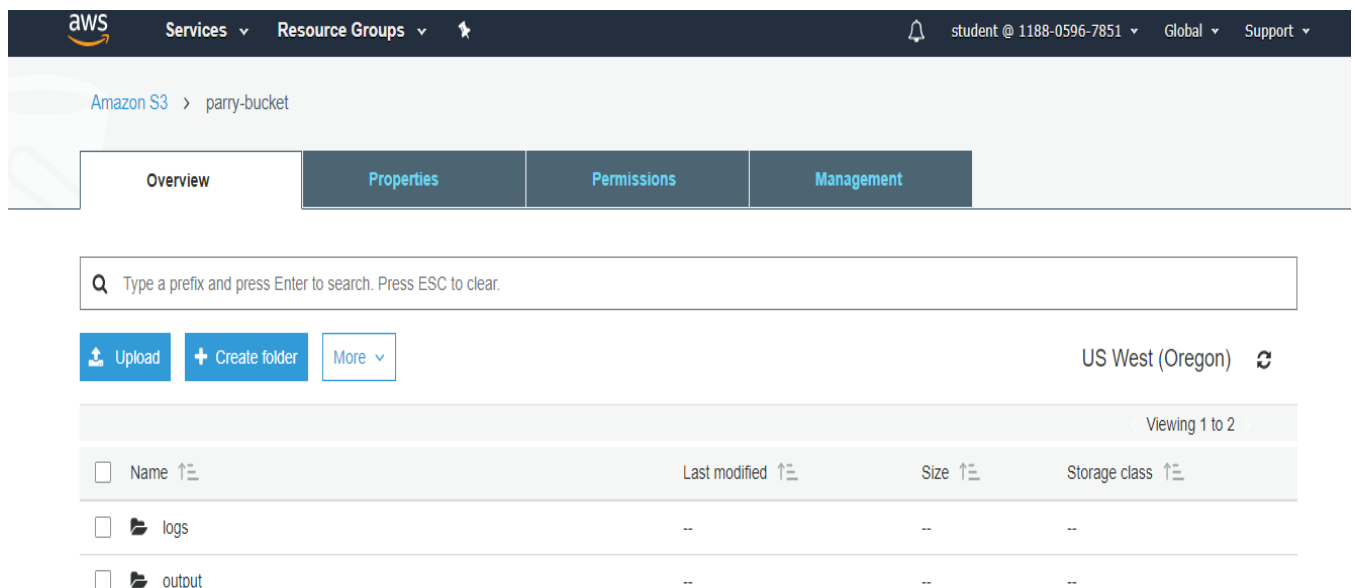


Fig 5 create S3 bucket under name parry-bucket

In this S3 bucket I created 2 folders name under logs, output, input. S3 is aws storage service providing

Services > storage > S3

Create bucket & then create folder

Create folder name

- Logs
- Input

Now I creating EMR cluster in EMR cluster there are major for configuration software, general , hardware , security configuration as shown in fig

aws Services Resource Groups

Create Cluster - Advanced Options [Go to quick options](#)

Step 1: Software and Steps
Step 2: Hardware
Step 3: General Cluster Settings
Step 4: Security

Software Configuration

Release **emr-5.13.0**

☒ Hadoop 2.8.3 ☐ Zeppelin 0.7.3 ☐ Livy 0.4.0
☐ Tez 0.8.4 ☐ Flink 1.4.0 ☐ Ganglia 3.7.2
☐ HBase 1.4.2 ☒ Pig 0.17.0 ☒ Hive 2.3.2
☐ Presto 0.194 ☐ ZooKeeper 3.4.10 ☐ MXNet 1.0.0
☐ Sqoop 1.4.6 ☐ Mahout 0.13.0 ☒ Hue 4.1.0
☐ Phoenix 4.13.0 ☐ Oozie 4.3.0 ☐ Spark 2.3.0
☐ HCatalog 2.3.2

AWS Glue Data Catalog settings (optional)

☐ Use for Hive table metadata

Edit software settings

☒ Enter configuration ☐ Load JSON from S3

`classification=config-file-name,properties=[myKey1=myValue1,myKey2=myValue2]`

Fig 6 creating cluster and e.m.r-5.6 as per requirement

Create Cluster - Quick Options [Go to advanced options](#)

Create Cluster - Quick Options [Go to advanced options](#)

General Configuration

Cluster name **My cluster**

☒ Logging

S3 folder **s3://aws-logs-000876196005-us-west-2/elasticmaprec**

Launch mode ☒ Cluster ☐ Step execution

Software configuration

Release **emr-5.13.0**

Applications

☒ Core Hadoop: Hadoop 2.8.3 with Ganglia 3.7.2, Hive 2.3.2, Hue 4.1.0, Mahout 0.13.0, Pig 0.17.0, and Tez 0.8.4

☐ HBase: HBase 1.4.2 with Ganglia 3.7.2, Hadoop 2.8.3, Hive 2.3.2, Hue 4.1.0, Phoenix 4.13.0, and ZooKeeper 3.4.10

☐ Presto: Presto 0.194 with Hadoop 2.8.3 HDFS and Hive 2.3.2 Metastore

☐ Spark: Spark 2.3.0 on Hadoop 2.8.3 YARN with Ganglia 3.7.2 and Zeppelin 0.7.3

☐ Use AWS Glue Data Catalog for table metadata

Fig 7 configuration for creating cluster

In software configuration I select EMR 5.6.0 which including Hadoop , pig, hive after clicking next few minute later my cluster is ready now I am doing in this project of hive script that processes amazon cloud front(cf) log files . the hive script & the sample cf logs and details available on S3 public bucket

When we click Add Step then fill out the blanks with suitable answer when we finish the cluster we need to wait for minute to complete the setup

I am running Hive on aws cluster & hive script running sample I clicked on a cluster list left hand side shown in a figure after that I clicked add step & fill hive program like name , output file etc, select the running cluster name & the summary information shows the cluster running and In the end in this hive lab initially I set up the cluster as per requirement the whole processes running on EMR the code I used was hive script provide by amazon, along with example of cloudfront logs.

The screenshot displays the AWS EMR console configuration page, divided into three main sections: Software configuration, Hardware configuration, and Security and access.

- Software configuration:** The 'Release' dropdown is set to 'emr-5.13.0'. Under 'Applications', 'Core Hadoop' is selected, which includes Hadoop 2.8.3, Ganglia 3.7.2, Hive 2.3.2, Hue 4.1.0, Mahout 0.13.0, Pig 0.17.0, and Tez 0.8.4. Other options like HBase, Presto, and Spark are unselected. The checkbox for 'Use AWS Glue Data Catalog for table metadata' is also unchecked.
- Hardware configuration:** The 'Instance type' is set to 'm3.xlarge'. The 'Number of instances' is set to '3', with a note indicating '(1 master and 2 core nodes)'.
- Security and access:** The 'EC2 key pair' dropdown shows 'Choose an option'. 'Permissions' are set to 'Default'. The 'EMR role' is 'EMR_DefaultRole' and the 'EC2 instance profile' is 'EMR_EC2_DefaultRole'. A link 'Learn how to create an EC2 key pair' is visible.

Fig 8 selecting hardware & security configuration

The screenshot shows the Amazon EMR console interface. The top navigation bar includes the AWS logo, 'Services', 'Resource Groups', and user information. The left sidebar lists navigation options: Amazon EMR, Clusters, Security configurations, VPC subnets, Events, and Help. The main content area displays the 'Cluster: CA cluster' in a 'Starting' state. It features several tabs: Summary, Monitoring, Hardware, Events, Steps, Configurations, and Bootstrap actions. The 'Summary' tab is active, showing details such as ID (j-1DOFETVQELJU4), Creation date (2018-05-08 13:57 UTC+1), Elapsed time (28 seconds), and Auto-terminate status (No). It also lists configuration details like Release label (emr-5.6.0), Hadoop distribution (Amazon 2.7.3), and Applications (Ganglia 3.7.2, Hive 2.1.1, Hue 3.12.0, Mahout 0.13.0, Pig 0.16.0, Tez 0.8.4). Network and hardware details include Availability zone, Subnet ID, and instance types (m3.xlarge). Security and access details show Key name, EC2 instance profile, EMR role, and security groups.

Fig 10 CA cluster dashboard

The screenshot shows the Amazon EMR console interface with the 'Cluster: CA cluster' in a 'Running' state. The 'Steps' tab is active, displaying a list of steps. At the top, there are buttons for 'Clone', 'Terminate', and 'AWS CLI export'. Below the cluster name, there are buttons for 'Add step', 'Clone step', and 'Cancel step'. A filter dropdown is set to 'All steps', showing '2 steps (all loaded)'. The steps table has columns for ID, Name, Status, Start time (UTC+1), Elapsed time, and Log files. The first step, 'AWS Hive example to process CF logs' (ID: s-2E16PLJ26O8G1), is in a 'Running' state. The second step, 'Setup hadoop debugging' (ID: s-K2DBN3JHWVFN), is in a 'Completed' state.

ID	Name	Status	Start time (UTC+1)	Elapsed time	Log files
s-2E16PLJ26O8G1	AWS Hive example to process CF logs	Running	2018-05-08 14:19 (UTC+1)	57 seconds	View logs
s-K2DBN3JHWVFN	Setup hadoop debugging	Completed	2018-05-08 14:05 (UTC+1)	2 seconds	View logs

Fig 11 steps dashboard hive running successfully


```
Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j2.properties Async: false
OK
Time taken: 4.663 seconds
Query ID = hadoop_20180508131942_a9d76294-67cd-4a93-9854-17e2b9004eb2
Total jobs = 1
Launching Job 1 out of 1
Waiting for Tez session and AM to be ready...
```

Status: Running (Executing on YARN cluster with App id application_1525784420224_0001)

```
Map 1: -/-      Reducer 2: 0/1
Map 1: 0/1      Reducer 2: 0/1
Map 1: 0/1      Reducer 2: 0/1
Map 1: 0(+1)/1  Reducer 2: 0/1
Map 1: 0(+1)/1  Reducer 2: 0/1
Map 1: 0(+1)/1  Reducer 2: 0/1
Map 1: 0(+1)/1  Reducer 2: 0/1
Map 1: 0(+1)/1  Reducer 2: 0/1
Map 1: 0(+1)/1  Reducer 2: 0/1
Map 1: 0(+1)/1  Reducer 2: 0/1
Map 1: 1/1      Reducer 2: 0/1
Map 1: 1/1      Reducer 2: 0(+1)/1
Map 1: 1/1      Reducer 2: 1/1
```

Moving data to directory s3://parry-bucket/output/os_requests

```
OK
Time taken: 35.366 seconds
Command exiting with ret '0'
```

Fig 12 Result of hive script

The screenshot shows the AWS S3 console interface for a bucket named '000000_0'. The 'Properties' tab is selected, displaying the following details:

- Owner:** labs+aws50
- Last modified:** May 8, 2018 2:20:18 PM GMT+0100
- Etag:** c79c153e7d67a0d64b5a8cc3940c6f61
- Storage class:** Standard
- Server-side encryption:** None
- Size:** 60
- Link:** https://s3-us-west-2.amazonaws.com/parry-bucket/output/os_requests/000000_0

Below the details, there are buttons for 'Open', 'Download', 'Download as', 'Make public', and 'Copy path'. To the right, a Notepad window titled '000000_0 - Notepad' is open, displaying the following text:

```
File Edit Format View Help
Android|855
Linux|813
MacOS|8520S
X|799
Windows|883
|OS|794
```

Fig 13 output of hive script in s3 bucket

Output file generate a notepad file in S3 bucket

Stream data and map reduce using EMR

In this project I am using EMR services I create one dummy txt file for streaming data & MapReduce using aws although I did streaming data in Kafka, flume & MapReduce in my local machine during the class labs but using aws for Hadoop MapReduce is totally complete experience it reduce my time 60% to perform tasks. I save the .txt from my local machine to aws s3 bucket under folder name Input after that click add step & I select streaming program after filling up all the blanks like mapper, reducer, input & output as shown in fig

After waiting for few minutes log files generated there are 3 log files mainly stderr , syslog , stdout

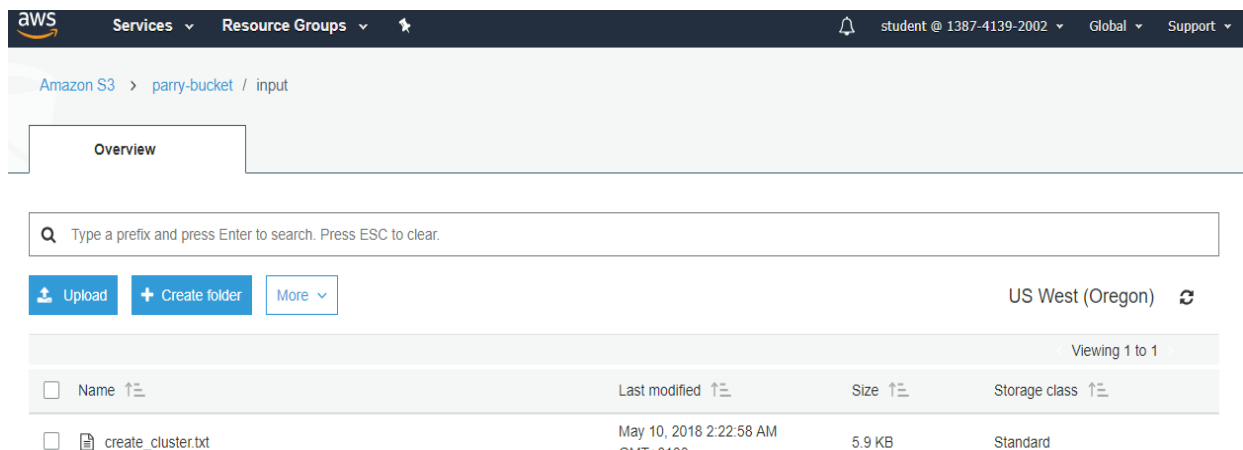


Fig 14 upload .txt file from local machine to S3

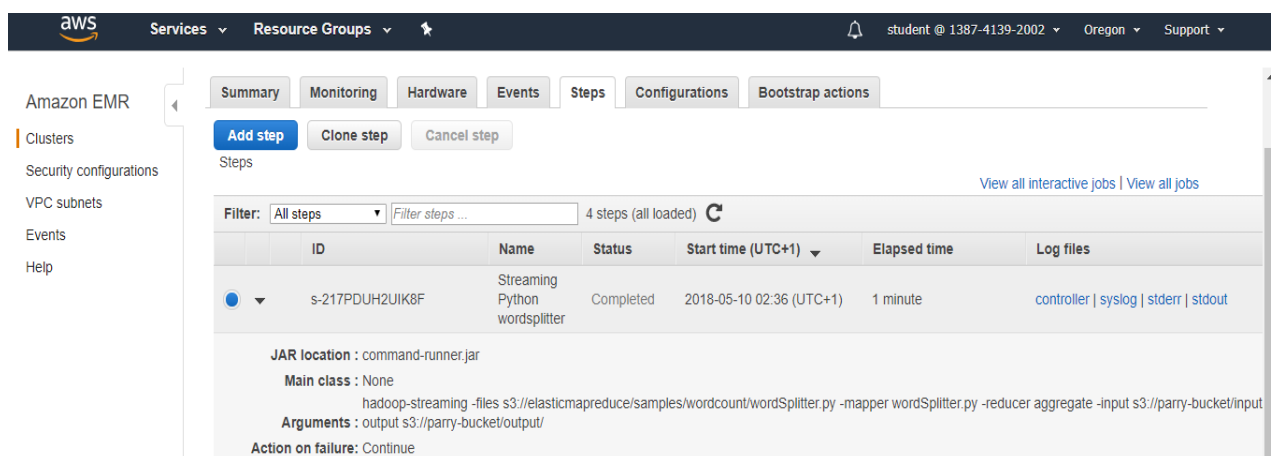


Fig 15 cluster status complete & logs file generated

aws

Services

Resource Groups

student @ 1387-4139-2002

Oregon

Support

Amazon EMR

Clusters

Security configurations

VPC subnets

Events

Help

Clone

Terminate

AWS CLI export

Cluster: CA cluster

Waiting

Cluster ready after last step completed.

Summary

Monitoring

Hardware

Events

Steps

Configurations

Bootstrap actions

Time	Event description	Source ID	Source type	Event type	Severity	Full date & time
May 10 02:38 AM	Amazon EMR cluster j-3R89RLEGVX0RE (CA cluster) finished running all pending steps at 2018-05-10 01:36 UTC.	j-3R89RLEGVX0RE	Cluster	Cluster State Change	INFO	May 10, 2018 at 02:38:37 AM (UTC+1)
May 10 02:38 AM	Step s-217PDUH2UIK8F (Streaming Python wordspli...) in Amazon EMR cluster j-3R89RLEGVX0RE (CA cluster) completed execution at 2018-05-10 01:38 UTC. The step started running at 2018-05-10 01:36 UTC and took 1 minutes to complete.	s-217PDUH2UIK8F	Step	Step State Change	INFO	May 10, 2018 at 02:38:36 AM (UTC+1)
May 10 02:36 AM	Amazon EMR cluster j-3R89RLEGVX0RE (CA cluster) began running steps at 2018-05-10 01:36 UTC.	j-3R89RLEGVX0RE	Cluster	Cluster State Change	INFO	May 10, 2018 at 02:36:52 AM (UTC+1)
May 10 02:36 AM	Step s-217PDUH2UIK8F (Streaming Python wordspli...) in Amazon EMR cluster j-3R89RLEGVX0RE (CA cluster) started running at 2018-05-10 01:36 UTC.	s-217PDUH2UIK8F	Step	Step State Change	INFO	May 10, 2018 at 02:36:51 AM (UTC+1)

Fig 16 list of events generated while creating cluster

Amazon S3

>

parry-bucket

/

output

Overview

Q

Type a prefix and press Enter to search. Press ESC to clear.

Upload

Create folder

More

US West (Oregon)

Viewing 1 to 4

<input type="checkbox"/>	Name	Last modified	Size	Storage class
<input type="checkbox"/>	_SUCCESS	May 10, 2018 2:38:02 AM GMT+0100	0 B	Standard
<input type="checkbox"/>	part-00000	May 10, 2018 2:37:56 AM GMT+0100	1.1 KB	Standard
<input type="checkbox"/>	part-00001	May 10, 2018 2:37:58 AM GMT+0100	933.0 B	Standard
<input type="checkbox"/>	part-00002	May 10, 2018 2:38:02 AM GMT+0100	1.0 KB	Standard

Fig 17 output file generated in S3 storage bucket

Testing



Fig 18 streaming data in 3 different file generated by EMR & stored in S3 parry-bucket

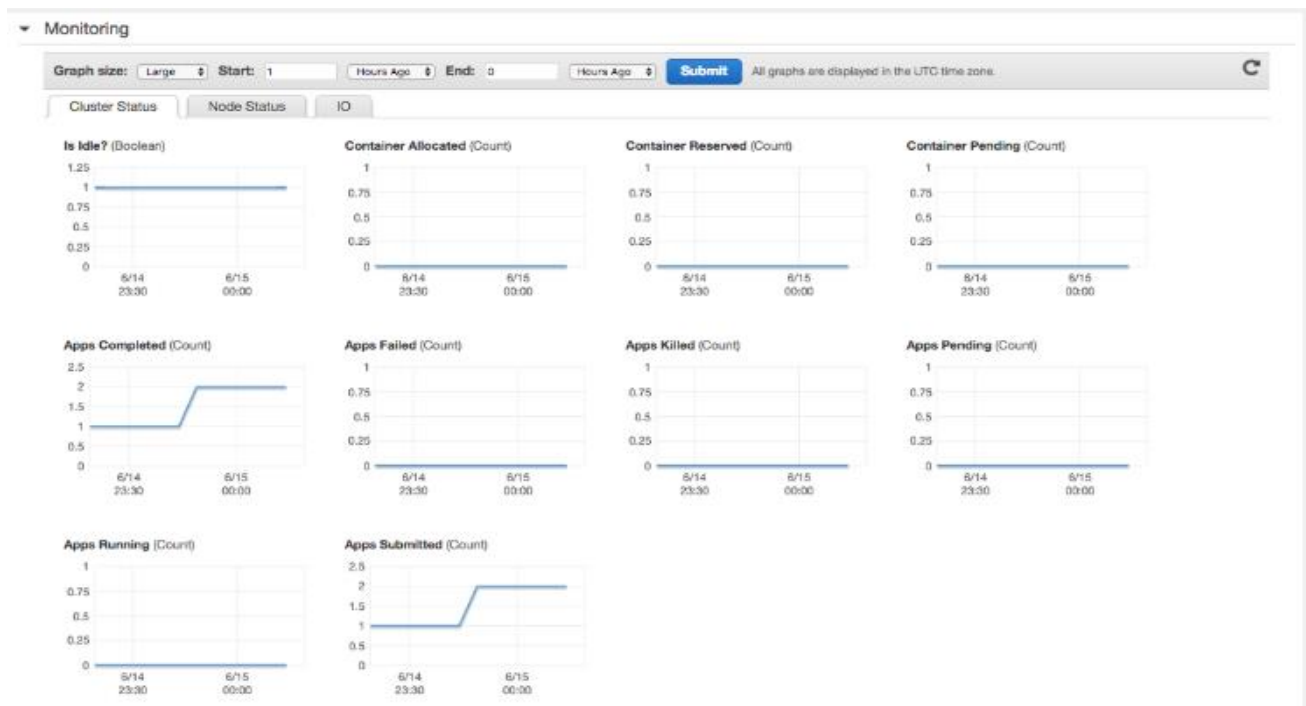


Fig 19 monitoring dashboard

here we can see all the activities we are doing in emr cluster like in above fig we can see that apps completed & apps submitted having sigmoid graph which means that our map-reduced successfully we created

After waiting for few minutes log files generated there are 3 log files mainly stderr , syslog , stout we can see the MapReduce

```
2018-05-10T01:36:51.080Z INFO Ensure step 4 jar file command-runner.jar
2018-05-10T01:36:51.080Z INFO StepRunner: Created Runner for step 4
INFO startExec 'hadoop jar /var/lib/aws/emr/step-runner/hadoop-jars/command-runner.jar hadoop-streaming -files s3://elasticmapreduce/samples/wordcount/wordSplitter.py -mapper wordSplitter.py -reducer aggregate -input s3://parry-bucket/input/create_cluster.txt -output s3://parry-bucket/output/'
INFO Environment:
PATH=/sbin:/usr/sbin:/bin:/usr/bin:/usr/local/sbin:/opt/aws/bin
LESS_TERMCAP_md=[01;38;5;208m
LESS_TERMCAP_me=[0m
HISTCONTROL=ignoredups
LESS_TERMCAP_mb=[01;31m
AWS_AUTO_SCALING_HOME=/opt/aws/apitools/as
UPSTART_JOB=rc
LESS_TERMCAP_se=[0m
HISTSIZE=1000
HADOOP_ROOT_LOGGER=INFO,DRFA
JAVA_HOME=/etc/alternatives/jre
AWS_DEFAULT_REGION=us-west-2
AWS_ELB_HOME=/opt/aws/apitools/elb
LESS_TERMCAP_us=[04;38;5;111m
EC2_HOME=/opt/aws/apitools/ec2
TERM=linux
XFILESEARCHPATH=/usr/dt/app-defaults/%L/Dt
runlevel=3
LANG=en_US.UTF-8
AWS_CLOUDWATCH_HOME=/opt/aws/apitools/mon
MAIL=/var/spool/mail/hadoop
LESS_TERMCAP_ue=[0m
LOGNAME=hadoop
PWD=/
LANGSH_SOURCED=1
HADOOP_CLIENT_OPTS=-Djava.io.tmpdir=/mnt/var/lib/hadoop/steps/s-217PDUH2UIK8F/tmp
_=/etc/alternatives/jre/bin/java
CONSOLETYPE=serial
RUNLEVEL=3
LESSOPEN=||/usr/bin/lesspipe.sh %s
previous=N
UPSTART_EVENTS=runlevel
AWS_PATH=/opt/aws
```

```

AWS_CLOUDWATCH_HOME=/opt/aws/apitools/mon
MAIL=/var/spool/mail/hadoop
LESS_TERMCAP_ue=[0m
LOGNAME=hadoop
PWD=/
LANGSH_SOURCED=1
HADOOP_CLIENT_OPTS=-Djava.io.tmpdir=/mnt/var/lib/hadoop/steps/s-217PDUH2UIK8F/tmp
_=/etc/alternatives/jre/bin/java
CONSOLETYPE=serial
RUNLEVEL=3
LESSOPEN=||/usr/bin/lesspipe.sh %s
previous=N
UPSTART_EVENTS=runlevel
AWS_PATH=/opt/aws
USER=hadoop
UPSTART_INSTANCE=
PREVLEVEL=N
HADOOP_LOGFILE=syslog
PYTHON_INSTALL_LAYOUT=amzn
HOSTNAME=ip-172-31-21-214
NLSPATH=/usr/dt/lib/nls/msg/%N.cat
HADOOP_LOG_DIR=/mnt/var/log/hadoop/steps/s-217PDUH2UIK8F
EC2_AMITOOL_HOME=/opt/aws/amitools/ec2
SHLVL=5
HOME=/home/hadoop
HADOOP_IDENT_STRING=hadoop
INFO redirectOutput to /mnt/var/log/hadoop/steps/s-217PDUH2UIK8F/stdout
INFO redirectError to /mnt/var/log/hadoop/steps/s-217PDUH2UIK8F/stderr
INFO Working dir /mnt/var/lib/hadoop/steps/s-217PDUH2UIK8F
INFO ProcessRunner started child process 28007 :
hadoop 28007 3344 1 01:36 ? 00:00:00 bash /usr/lib/hadoop/bin/hadoop jar /var/lib/aws/emr/step-runner/hadoop-jars/command-runner.jar hadoop-streaming -files
s3://elasticmapreduce/samples/wordcount/wordSplitter.py -mapper wordSplitter.py -reducer aggregate -input s3://parry-bucket/input/create_cluster.txt -output s3://parry-bucket/output/
2018-05-10T01:36:55.092Z INFO HadoopJarStepRunner.Runner: startRun() called for s-217PDUH2UIK8F Child Pid: 28007
INFO Synchronously wait child process to complete : hadoop jar /var/lib/aws/emr/step-runner/hadoop-...
INFO waitProcessCompletion ended with exit code 0 : hadoop jar /var/lib/aws/emr/step-runner/hadoop-...
INFO total process run time: 72 seconds
2018-05-10T01:38:05.321Z INFO Step created jobs: job_1525913736398_0002
2018-05-10T01:38:05.322Z INFO Step succeeded with exitCode 0 and took 72 seconds

```

Activate Windows

Fig 20 log file created streaming data

```

2018-05-10 01:36:59,936 INFO org.apache.hadoop.yarn.client.api.impl.TimelineClientImpl (main): Timeline service address: http://ip-172-31-21-214.us-west-2.compute.internal:8188/ws/v1/timeline/
2018-05-10 01:36:59,949 INFO org.apache.hadoop.yarn.client.RMProxy (main): Connecting to ResourceManager at ip-172-31-21-214.us-west-2.compute.internal/172.31.21.214:8032
2018-05-10 01:37:00,348 INFO org.apache.hadoop.yarn.client.api.impl.TimelineClientImpl (main): Timeline service address: http://ip-172-31-21-214.us-west-2.compute.internal:8188/ws/v1/timeline/
2018-05-10 01:37:00,348 INFO org.apache.hadoop.yarn.client.RMProxy (main): Connecting to ResourceManager at ip-172-31-21-214.us-west-2.compute.internal/172.31.21.214:8032
2018-05-10 01:37:00,787 INFO com.amazon.ws.emr.hadoop.fs.s3n.S3NativeFileSystem (main): Opening 's3://elasticmapreduce/samples/wordcount/wordSplitter.py' for reading
2018-05-10 01:37:01,611 INFO com.hadoop.compression.lzo.GPLNativeCodeLoader (main): Loaded native gpl library
2018-05-10 01:37:01,613 INFO com.hadoop.compression.lzo.LzoCodec (main): Successfully loaded & initialized native-lzo library [hadoop-lzo rev 154f1ef53e2d6ed126b0957d7995e0a610947608]
2018-05-10 01:37:01,629 INFO org.apache.hadoop.mapred.FileInputFormat (main): Total input paths to process : 1
2018-05-10 01:37:01,683 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): number of splits:8
2018-05-10 01:37:01,832 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): Submitting tokens for job: job_1525913736398_0002
2018-05-10 01:37:02,100 INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl (main): Submitted application application_1525913736398_0002
2018-05-10 01:37:02,218 INFO org.apache.hadoop.mapreduce.Job (main): The url to track the job: http://ip-172-31-21-214.us-west-2.compute.internal:20888/proxy/application_1525913736398_0002/
2018-05-10 01:37:02,220 INFO org.apache.hadoop.mapreduce.Job (main): Running job: job_1525913736398_0002
2018-05-10 01:37:09,374 INFO org.apache.hadoop.mapreduce.Job (main): Job job_1525913736398_0002 running in uber mode : false
2018-05-10 01:37:09,376 INFO org.apache.hadoop.mapreduce.Job (main): map 0% reduce 0%
2018-05-10 01:37:31,554 INFO org.apache.hadoop.mapreduce.Job (main): map 13% reduce 0%
2018-05-10 01:37:35,578 INFO org.apache.hadoop.mapreduce.Job (main): map 25% reduce 0%
2018-05-10 01:37:39,602 INFO org.apache.hadoop.mapreduce.Job (main): map 38% reduce 0%
2018-05-10 01:37:41,614 INFO org.apache.hadoop.mapreduce.Job (main): map 50% reduce 0%
2018-05-10 01:37:42,620 INFO org.apache.hadoop.mapreduce.Job (main): map 63% reduce 0%
2018-05-10 01:37:44,636 INFO org.apache.hadoop.mapreduce.Job (main): map 75% reduce 0%
2018-05-10 01:37:52,689 INFO org.apache.hadoop.mapreduce.Job (main): map 88% reduce 0%
2018-05-10 01:37:54,703 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 0%
2018-05-10 01:37:57,720 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 33%
2018-05-10 01:37:59,737 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 67%
2018-05-10 01:38:02,755 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 100%
2018-05-10 01:38:03,768 INFO org.apache.hadoop.mapreduce.Job (main): Job job_1525913736398_0002 completed successfully
2018-05-10 01:38:03,890 INFO org.apache.hadoop.mapreduce.Job (main): Counters: 56
File System Counters
FILE: Number of bytes read=4390
FILE: Number of bytes written=1442398
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=752

```

Activate Windows

```

    >>: Number of write operations=0
Job Counters
    Killed map tasks=1
    Killed reduce tasks=1
    Launched map tasks=8
    Launched reduce tasks=3
    Data-local map tasks=8
    Total time spent by all maps in occupied slots (ms)=8052660
    Total time spent by all reduces in occupied slots (ms)=3293280
    Total time spent by all map tasks (ms)=178948
    Total time spent by all reduce tasks (ms)=36592
    Total vcore-milliseconds taken by all map tasks=178948
    Total vcore-milliseconds taken by all reduce tasks=36592
    Total megabyte-milliseconds taken by all map tasks=257685120
    Total megabyte-milliseconds taken by all reduce tasks=105384960
Map-Reduce Framework
    Map input records=50
    Map output records=971
    Map output bytes=20259
    Map output materialized bytes=7231
    Input split bytes=752
    Combine input records=971
    Combine output records=584
    Reduce input groups=342
    Reduce shuffle bytes=7231
    Reduce input records=584
    Reduce output records=342
    Spilled Records=1168
    Shuffled Maps =24
    Failed Shuffles=0
    Merged Map outputs=24
    GC time elapsed (ms)=4628
    CPU time spent (ms)=25760
    Physical memory (bytes) snapshot=5896880128
    Virtual memory (bytes) snapshot=39387328512
    Total committed heap usage (bytes)=5829033984
Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
File Input Format Counters
    Bytes Read=26996
File Output Format Counters
    Bytes Written=3121
2018-05-10 01:38:03,890 INFO org.apache.hadoop.streaming.StreamJob (main): Output directory: s3://parry-bucket/output/

```

Fig 21 showing map-reduce job and 100% mapping , reducing

Conclusion

In this project I did two labs one is hive script and another is streaming data , map-reducing job using EMR cluster I have completed both the task successfully using AWS really fast and productive in terms of time & money setup new Hadoop cluster is a hardcore pain during my class lab took me 5-6 hrs after going step by step but in AWS it's 10 min task step a Hadoop cluster and computational power is really big problem which AWS is solving like if I have to run 3-4 cluster at time in my machine I need 12+ gb ram with supporting processor also in case AWS I can scale up as much as I can or depending upon my requirement there are few limitation with AWS is charging higher cost for clusters compare to other cloud providers Dependency also issue with AWS like S3 storage went down last year and because of that companies like Pinterest also went down so depending upon only one service provider make things difficult

Reference

1. amazon. 2018. amazon. [ONLINE] Available at: <https://aws.amazon.com/what-is-cloud-computing/>. [Accessed 16 May 2018].
2. amazon. 2018. amazon. [ONLINE] Available at: <https://aws.amazon.com/documentation/emr/>. [Accessed 16 May 2018].
3. amazon. 2018. Zillow Provides Near-Real-Time Home-Value Estimates Using Amazon Kinesis. [ONLINE] Available at: https://d1.awsstatic.com/case-studies/PDF%20Case%20Studies/AWS-Casestudy_Zillow.pdf. [Accessed 21 May 2018].
4. Expedia. 2018. Expedia Increases Agility and Resiliency by Going All In on AWS. [ONLINE] Available at: <https://aws.amazon.com/solutions/case-studies/expedia/>. [Accessed 21 May 2018].
5. Yelp. 2018. Yelp Case Study. [ONLINE] Available at: <https://aws.amazon.com/solutions/case-studies/yelp/>. [Accessed 21 May 2018].